

Enhanced Regulatory Calculation Pipeline Architecture

Executive Summary

This document outlines the architecture for a high-performance regulatory calculation pipeline that processes large-scale financial data for liquidity calculations. The system features a sophisticated adjustment mechanism that allows rule-based modifications at any stage while maintaining complete data lineage and traceability.

Table of Contents

- 1. Current Architecture Overview
- 2. Core Components
- 3. Adjustment System Design
- 4. Incremental Recalculation Strategy
- 5. Performance Optimization
- 6. S3 to Redshift to PowerBI Optimization
- 7. Implementation Patterns
- 8. Monitoring and Observability

Current Architecture Overview

DAG-Based Pipeline Architecture




Diagram 0

Key Architectural Principles

- 1. **DAG-Based Workflow:** Each pipeline is an Airflow DAG with operators as nodes
- 2. **Lineage Preservation:** Every operator maintains row-level tracing through primary keys
- 3. **Smart Persistence:** Files written only when aggregations would lose tracing
- 4. **Flexible Adjustments:** Adjustment operators can be injected at any DAG node
- 5. **Visual Configuration:** DAG builder UI for pipeline construction

Core Components

Operator Architecture




Diagram 1

DAG Execution Flow

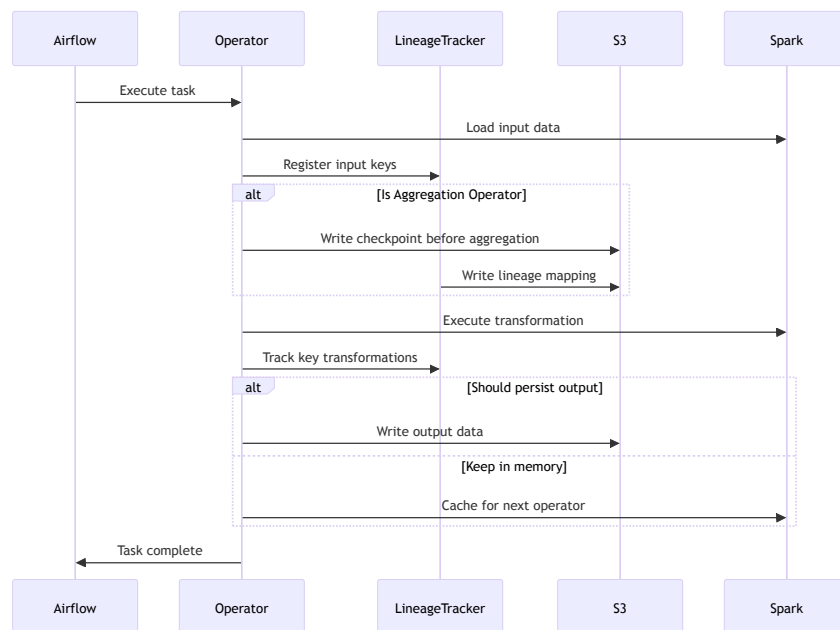


Diagram 2

Adjustment System Design

Flexible Adjustment Injection

```

class AdjustmentOperator(BaseOperator):
    """
    Can be injected at any point in the DAG to apply adjustments
    """

    def __init__(self, operator_id: str, adjustment_rules: DataFrame):
        self.operator_id = operator_id
        self.adjustment_rules = adjustment_rules
        self.maintains_keys = True # Always preserves lineage
  
```

```

def execute(self, input_df: DataFrame) -> DataFrame:
    # Apply adjustments while maintaining full lineage
    adjusted_df = self.apply_rule_based_adjustments(input_df)

    # Track what was adjusted
    adjustment_tracking = self.create_adjustment_audit(
        input_df,
        adjusted_df,
        self.adjustment_rules
    )

    # Store adjustment lineage
    self.persist_adjustment_lineage(adjustment_tracking)

    return adjusted_df

def supports_parallel_execution(self) -> bool:
    """Adjustments can run in parallel scenarios"""
    return True

class DAGBuilder:
    """
    Visual DAG builder that allows adjustment injection
    """

    def inject_adjustment_operator(
        self,
        dag: DAG,
        after_operator_id: str,
        adjustment_rules: DataFrame
    ) -> DAG:
        """
        Inject an adjustment operator at any point in the DAG
        """

        # Create new adjustment operator
        adj_op = AdjustmentOperator(
            f"adjustment_{after_operator_id}",
            adjustment_rules
        )

        # Rewire DAG to include adjustment
        original_downstream = dag.get_downstream(after_operator_id)
        dag.add_edge(after_operator_id, adj_op.operator_id)

        for downstream_op in original_downstream:
            dag.add_edge(adj_op.operator_id, downstream_op)

        return dag

```

Smart Persistence Strategy

```

class SmartPersistenceManager:
    """
    Determines when to persist data based on lineage requirements
    """

    def __init__(self):
        self.lineage_tracker = LineageTracker()

    def should_persist(self, operator: BaseOperator, input_df:
        DataFrame) -> bool:
        """
        Persist data when:
        1. Aggregation will lose row-level keys
        2. Explicit checkpoint requested
        3. Memory pressure threshold exceeded
        """

        if isinstance(operator, AggregationOperator):
            # Always persist before aggregations
            return True

        if operator.is_checkpoint:
            # Explicit checkpoint requested
            return True

        if self.estimate_memory_usage(input_df) >
            self.memory_threshold:
            # Memory pressure - persist to S3
            return True

        return False

    def persist_with_lineage(self, df: DataFrame, operator_id: str,
        path: str):
        """
        Persist data and its lineage information
        """

        # Write data
        df.write.mode("overwrite").parquet(path)

        # Write lineage metadata
        lineage_info = {
            "operator_id": operator_id,
            "row_count": df.count(),
            "key_columns": self.identify_key_columns(df),
            "schema": df.schema.json(),
            "timestamp": datetime.now().isoformat()
        }

        self.write_lineage_metadata(lineage_info, f"{path}/_lineage")

```

DAG-Based Pipeline Implementation

Visual DAG Builder

```

@dataclass
class OperatorConfig:
    """Configuration for each operator in the DAG"""
    operator_type: str
    operator_id: str
    parameters: Dict[str, Any]
    inputs: List[str]
    outputs: List[str]

class VisualDAGBuilder:
    """
    Allows users to visually build and configure DAGs
    """

    def __init__(self):
        self.operators = {}
        self.connections = []

    def add_operator(self, config: OperatorConfig) -> 'VisualDAGBuilder':
        """Add an operator to the DAG"""

        operator_class = self.get_operator_class(config.operator_type)
        operator = operator_class(**config.parameters)
        operator.operator_id = config.operator_id

        self.operators[config.operator_id] = operator
        return self

    def connect(self, from_id: str, to_id: str) -> 'VisualDAGBuilder':
        """Connect two operators"""
        self.connections.append((from_id, to_id))
        return self

    def inject_adjustment(
        self,
        after_operator_id: str,
        adjustment_rules: DataFrame
    ) -> 'VisualDAGBuilder':
        """Inject an adjustment operator after specified operator"""

        adj_id = f"adj_{after_operator_id}_{uuid.uuid4().hex[:8]}"
        adj_config = OperatorConfig(
            operator_type="adjustment",
            operator_id=adj_id,
            parameters={"rules": adjustment_rules},

```

```

        inputs=[after_operator_id],
        outputs=[]
    )

    self.add_operator(adj_config)

    # Rewire connections
    new_connections = []
    for from_id, to_id in self.connections:
        if from_id == after_operator_id:
            new_connections.append((from_id, adj_id))
            new_connections.append((adj_id, to_id))
        else:
            new_connections.append((from_id, to_id))

    self.connections = new_connections
    return self

def build_airflow_dag(self) -> DAG:
    """Convert visual DAG to Airflow DAG"""

    dag = DAG(
        'regulatory_calculation_pipeline',
        schedule_interval='@daily',
        catchup=False
    )

    # Create Airflow tasks
    tasks = {}
    for op_id, operator in self.operators.items():
        task = PythonOperator(
            task_id=op_id,
            python_callable=operator.execute_with_lineage,
            dag=dag
        )
        tasks[op_id] = task

    # Set dependencies
    for from_id, to_id in self.connections:
        tasks[from_id] >> tasks[to_id]

    return dag

```

Example DAG Configuration

```

class RegulatoryCalculationDAG:
    """
    Example of a complete regulatory calculation DAG
    """

    def build_dag(self) -> DAG:

```

```
builder = VisualDAGBuilder()

# Data ingestion
builder.add_operator(OperatorConfig(
    operator_type="ingestion",
    operator_id="ingest_deals",
    parameters={"source": "s3://data/raw/deals"},
    inputs=[],
    outputs=["deals_df"]
))

# Enrichment
builder.add_operator(OperatorConfig(
    operator_type="enrichment",
    operator_id="enrich_deals",
    parameters={"reference_data": "s3://data/reference"},
    inputs=["deals_df"],
    outputs=["enriched_df"]
))

# Liquidity calculation
builder.add_operator(OperatorConfig(
    operator_type="calculation",
    operator_id="calc_liquidity",
    parameters={"calculation_type": "lcr"},
    inputs=["enriched_df"],
    outputs=["liquidity_df"]
))

# Aggregation (will trigger persistence)
builder.add_operator(OperatorConfig(
    operator_type="aggregation",
    operator_id="agg_by_entity",
    parameters={"group_by": ["entity_id", "date"]},
    inputs=["liquidity_df"],
    outputs=["aggregated_df"]
))

# Star schema generation
builder.add_operator(OperatorConfig(
    operator_type="star_schema",
    operator_id="generate_star",
    parameters={
        "fact_table": "liquidity_metrics",
        "dimensions": ["entity", "date", "product"]
    },
    inputs=["aggregated_df"],
    outputs=["star_schema"]
))

# Connect operators
```

```

(builder
    .connect("ingest_deals", "enrich_deals")
    .connect("enrich_deals", "calc_liquidity")
    .connect("calc_liquidity", "agg_by_entity")
    .connect("agg_by_entity", "generate_star"))

return builder.build_airflow_dag()

```

Parallel Execution for Impact Analysis

```

class ParallelAdjustmentAnalyzer:
    """
    Run multiple adjustment scenarios in parallel to analyze impact
    """

    def __init__(self, base_dag: DAG):
        self.base_dag = base_dag

    def analyze_adjustment_impact(
        self,
        adjustment_scenarios: List[AdjustmentScenario],
        injection_points: List[str]
    ) -> DataFrame:
        """
        Run parallel DAGs with different adjustments
        """

        results = []

        # Create parallel DAGs for each scenario
        parallel_dags = []
        for scenario in adjustment_scenarios:
            for injection_point in injection_points:
                # Clone base DAG
                scenario_dag = self.clone_dag(self.base_dag)

                # Inject adjustment
                scenario_dag = self.inject_adjustment(
                    scenario_dag,
                    injection_point,
                    scenario.rules
                )

                parallel_dags.append({
                    "dag": scenario_dag,
                    "scenario": scenario.name,
                    "injection_point": injection_point
                })

        # Execute in parallel
        with ThreadPoolExecutor(max_workers=10) as executor:

```



```

futures = []
for dag_info in parallel_dags:
    future = executor.submit(
        self.execute_dag,
        dag_info["dag"],
        dag_info["scenario"],
        dag_info["injection_point"]
    )
    futures.append(future)

# Collect results
for future in as_completed(futures):
    result = future.result()
    results.append(result)

# Compare results
return self.create_impact_analysis(results)

def create_impact_analysis(self, results: List[Dict]) ->
    DataFrame:
    """
    Create comprehensive impact analysis report
    """

    analysis_df = self.spark.createDataFrame(results)

    # Calculate impact metrics
    impact_metrics = analysis_df.groupBy("scenario",
        "injection_point").agg(
        F.sum("records_affected").alias("total_records_affected"),

        F.avg("liquidity_change_pct").alias("avg_liquidity_change"),
        F.max("max_entity_impact").alias("max_entity_impact"),

        F.collect_list("affected_entities").alias("all_affected_entities")
    )

    return impact_metrics

```

Adjustment System Design

Adjustment Rule Structure

```

adjustment_rule = {
    "rule_id": "ADJ_2024_001",
    "stage": "deal_enrichment", # or "aggregation", "final_calc"
    "conditions": [
        {"field": "product_type", "operator": "in", "value": ["BOND",
            "LOAN"]},
        {"field": "maturity_date", "operator": ">", "value": "2025-01-
            01"}
    ]
}

```

```

    ],
    "adjustments": [
      {"field": "liquidity_factor", "operation": "multiply",
       "value": 0.95},
      {"field": "haircut", "operation": "add", "value": 0.05}
    ],
    "effective_date": "2024-01-15",
    "expiry_date": "2024-12-31",
    "approval": {"user": "risk_admin", "timestamp": "2024-01-14T10:00:00Z"}
  }
}

```

Adjustment Application Flow

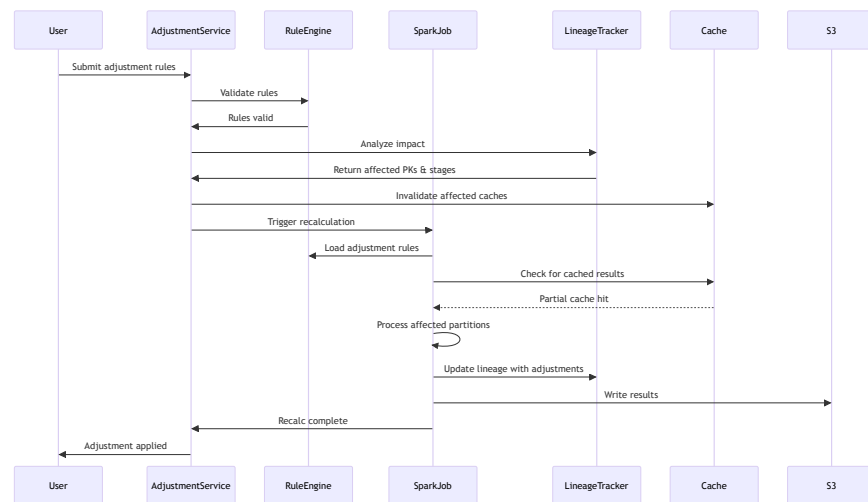


Diagram 3

Incremental Recalculation Strategy

Dependency Graph Management

```

class DependencyGraph:
    def __init__(self):
        self.graph = nx.DiGraph()
        self.stage_cache = {}

    def add_stage_dependency(self, source_stage, target_stage,
                             pk_mapping):
        self.graph.add_edge(source_stage, target_stage,
                             pk_mapping=pk_mapping)

    def find_minimal_recalc_path(self, adjustment_rule):
        affected_stage = adjustment_rule['stage']
        affected_pks = self.get_affected_pks(adjustment_rule)

        # Find all downstream stages

```

```

downstream_stages = nx.descendants(self.graph, affected_stage)

# Build minimal recalc plan
recalc_plan = RecalcPlan()
for stage in nx.topological_sort(self.graph):
    if stage in downstream_stages or stage == affected_stage:
        stage_pks = self.propagate_pks(affected_pks,
                                       affected_stage, stage)
        recalc_plan.add_stage(stage, stage_pks)

return recalc_plan

```

Incremental Processing Pattern

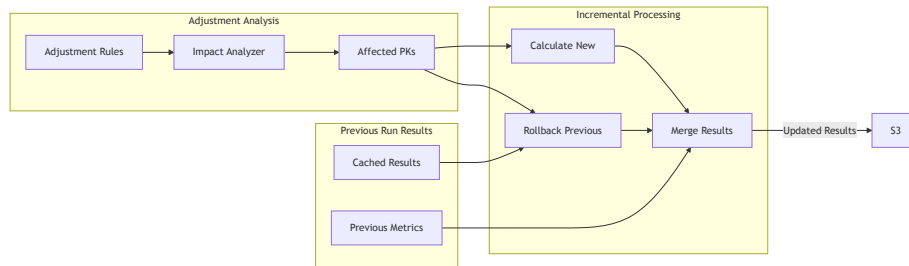


Diagram 4

Performance Optimization

Multi-Layer Caching Strategy

```

class CacheManager:
    def __init__(self, s3_bucket, redis_cluster):
        self.s3_cache = S3Cache(s3_bucket)
        self.memory_cache = SparkCacheManager()
        self.redis_cache = RedisCache(redis_cluster)

    def get_or_compute(self, stage_id, partition_key, compute_func):
        # L1: In-memory Spark cache
        cached = self.memory_cache.get(stage_id, partition_key)
        if cached:
            return cached

        # L2: Redis cache for hot data
        cached = self.redis_cache.get(stage_id, partition_key)
        if cached:
            df = self.deserialize_from_redis(cached)
            self.memory_cache.put(stage_id, partition_key, df)
            return df

        # L3: S3 cache for cold data
        cached = self.s3_cache.get(stage_id, partition_key)

```

```

if cached:
    df = spark.read.parquet(cached)
    self.memory_cache.put(stage_id, partition_key, df)
    return df

# Compute and cache at all levels
result = compute_func()
self.cache_result(stage_id, partition_key, result)
return result

```

Parallel Execution Strategy

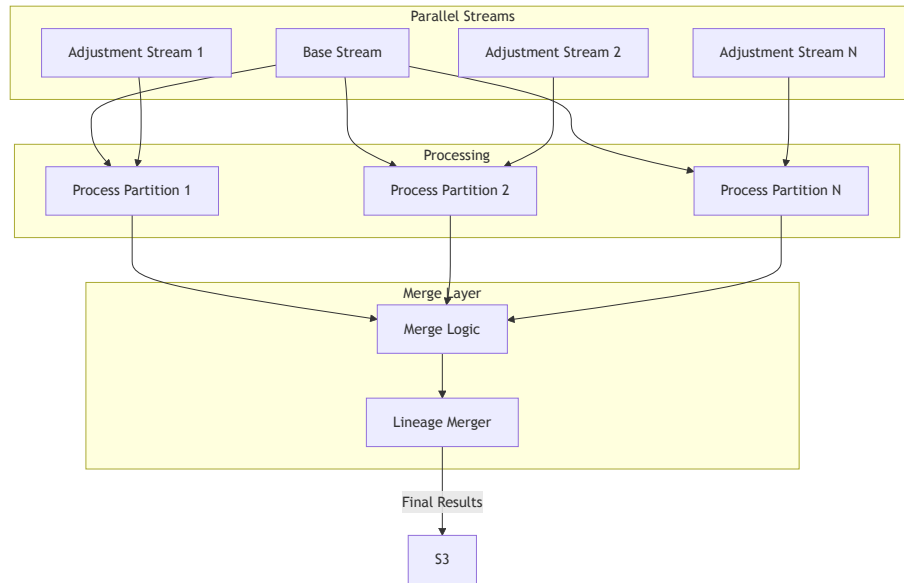


Diagram 5

S3 to Redshift to PowerBI Optimization

Direct Write vs. S3 Staging Approach

Architecture Comparison



Diagram 6

Direct Write Implementation

Direct Write Implementation

```

# Direct Redshift Writer with Star Schema Transformation
class RedshiftStarSchemaWriter:

```

```

def __init__(self, spark_session, jdbc_url, temp_s3_path):
    self.spark = spark_session
    self.jdbc_url = jdbc_url
    self.temp_s3_path = temp_s3_path
    self.redshift_props = {
        "user": "admin",
        "password": "password",
        "driver": "com.amazon.redshift.jdbc42.Driver",
        "tempdir": temp_s3_path,
        "aws_iam_role": "arn:aws:iam::account:role/RedshiftRole"
    }

def transform_to_star_schema(self, raw_df):
    """Transform raw data into star schema format"""

    # Create dimension tables
    dim_entity = self.create_entity_dimension(raw_df)
    dim_date = self.create_date_dimension(raw_df)
    dim_product = self.create_product_dimension(raw_df)
    dim_adjustment = self.create_adjustment_dimension(raw_df)

    # Create fact table with foreign keys
    fact_liquidity = self.create_fact_table(
        raw_df, dim_entity, dim_date, dim_product, dim_adjustment
    )

    return {
        "dim_entity": dim_entity,
        "dim_date": dim_date,
        "dim_product": dim_product,
        "dim_adjustment": dim_adjustment,
        "fact_liquidity": fact_liquidity
    }

def write_with_performance_optimization(self, tables_dict):
    """Write to Redshift with various optimization strategies"""

    for table_name, df in tables_dict.items():
        if table_name.startswith("dim_"):
            # Dimension tables – use UPSERT pattern
            self.upsert_dimension(table_name, df)
        else:
            # Fact tables – use batch insert or COPY based on size
            if df.count() < 100000: # Threshold for direct write
                self.direct_write(table_name, df)
            else:
                self.staged_write(table_name, df)

def direct_write(self, table_name, df):
    """Direct JDBC write for smaller datasets"""

```

```

# Optimize for direct write
optimized_df = (df
    .coalesce(10) # Reduce parallelism for small data
    .sortWithinPartitions("entity_key", "date_key")
)

# Write with batch optimization
(optimized_df.write
    .mode("append")
    .option("batchsize", 10000)
    .option("truncate", "false")
    .jdbc(self.jdbc_url, table_name,
properties=self.redshift_props)
)

def staged_write(self, table_name, df):
    """S3-staged write for larger datasets"""

    # Write to S3 first
    temp_path = f"{self.temp_s3_path}/{table_name}/{uuid.uuid4()}"

    (df.repartition(self.calculate_optimal_partitions(df))
        .write
        .mode("overwrite")
        .parquet(temp_path))

    # Use COPY command
    copy_command = f"""
COPY {table_name}
FROM '{temp_path}'
IAM_ROLE '{self.redshift_props["aws_iam_role"]}'
FORMAT AS PARQUET
COMPUPDATE PRESET
STATUPDATE ON;
"""

    self.execute_redshift_command(copy_command)

    # Clean up S3
    self.cleanup_s3(temp_path)

def upsert_dimension(self, table_name, df):
    """Upsert pattern for dimension tables"""

    # Write to staging table first
    staging_table = f"{table_name}_staging"

    (df.write
        .mode("overwrite")
        .jdbc(self.jdbc_url, staging_table,
properties=self.redshift_props)
    )

```

```

# Perform MERGE operation
merge_sql = f"""
BEGIN TRANSACTION;

-- Update existing records
UPDATE {table_name}
SET
    last_updated = s.last_updated,
    is_active = s.is_active,
    attributes = s.attributes
FROM {staging_table} s
WHERE {table_name}.natural_key = s.natural_key;

-- Insert new records
INSERT INTO {table_name}
SELECT s.* FROM {staging_table} s
LEFT JOIN {table_name} t ON s.natural_key = t.natural_key
WHERE t.natural_key IS NULL;

DROP TABLE {staging_table};

COMMIT;
"""

self.execute_redshift_command(merge_sql)

```

Performance Comparison & Trade-offs

Approach	Pros	Cons	Best Use Case
S3 Staging (Current)	<ul style="list-style-type: none"> • Highly scalable • Fault tolerant • Best for large batches • S3 provides backup 	<ul style="list-style-type: none"> • Higher latency • Two-step process • S3 costs 	Regular batch loads >1M records
Direct Write	<ul style="list-style-type: none"> • Lower latency • Real-time updates • No S3 costs • Immediate availability 	<ul style="list-style-type: none"> • Connection overhead • Less fault tolerant • Poor for large batches • Redshift 	Small updates, adjustments <100K records

Approach	Pros	Cons	Best Use Case
		load impact	
Hybrid (Recommended)	<ul style="list-style-type: none"> • Optimal performance • Flexible approach • Cost effective • Handles all scenarios 	<ul style="list-style-type: none"> • More complex logic • Requires monitoring • Multiple code paths 	Production systems with mixed workloads

Optimized Hybrid Architecture Implementation

```

class HybridRedshiftWriter:
    def __init__(self, spark_session, config):
        self.spark = spark_session
        self.config = config
        self.metrics = MetricsCollector()

    def write_to_redshift(self, df, table_name, write_mode="auto"):
        """Intelligent routing based on data characteristics"""

        # Analyze data characteristics
        row_count = df.count()
        is_adjustment = "adjustment_id" in df.columns
        partition_count = df.rdd.getNumPartitions()

        # Decision logic
        if write_mode == "auto":
            if is_adjustment and row_count < 50000:
                # Adjustments need immediate visibility
                return self.direct_upsert(df, table_name)
            elif row_count < 100000:
                # Small batches - direct write
                return self.optimized_direct_write(df, table_name)
            else:
                # Large batches - S3 staging
                return self.staged_copy_write(df, table_name)

    def optimized_direct_write(self, df, table_name):
        """Direct write with connection pooling and batching"""

        start_time = time.time()

        try:
            # Transform to star schema if needed
            if self.requires_transformation(table_name):
                df = self.transform_to_star_schema(df)

```



```

# Optimize DataFrame for direct write
optimized_df = self.optimize_for_direct_write(df)

# Use connection pooling
with self.get_connection_pool() as conn_pool:
    (optimized_df.write
     .mode("append")
     .option("batchsize", 25000)
     .option("numPartitions", 4) # Limit concurrent
connections
     .option("isolationLevel", "READ_UNCOMMITTED")
     .jdbc(self.config.jdbc_url, table_name,
           connection_properties=conn_pool.properties))

self.metrics.record_write(
    method="direct",
    table=table_name,
    records=df.count(),
    duration=time.time() - start_time
)

except Exception as e:
    # Fallback to S3 staging on failure
    self.logger.warning(f"Direct write failed, falling back to
S3: {e}")
    return self.staged_copy_write(df, table_name)

def optimize_for_direct_write(self, df):
    """Optimize DataFrame for JDBC write performance"""

    return (df
            # Reduce partitions to limit connections
            .coalesce(min(df.rdd.getNumPartitions(), 10))
            # Sort for better compression
            .sortWithinPartitions("date_key", "entity_key")
            # Cache if we'll read multiple times
            .cache()
            )

```

Star Schema Transformation in Spark

```

class StarSchemaTransformer:
    def __init__(self, spark_session):
        self.spark = spark_session

    def create_date_dimension(self, df):
        """Create date dimension with proper attributes"""

        dates_df = df.select("calculation_date").distinct()

```

```

return dates_df.select(
    F.monotonically_increasing_id().alias("date_key"),
    F.col("calculation_date"),
    F.year("calculation_date").alias("year"),
    F.quarter("calculation_date").alias("quarter"),
    F.month("calculation_date").alias("month"),
    F.dayofmonth("calculation_date").alias("day"),
    F.dayofweek("calculation_date").alias("day_of_week"),
    F.weekofyear("calculation_date").alias("week_of_year"),
    F.when(F.dayofweek("calculation_date").isin([1,7]), True)
      .otherwise(False).alias("is_weekend"),
    F.last_day("calculation_date").alias("month_end_date")
).distinct()

def create_fact_table(self, raw_df, dim_entity, dim_date,
dim_product):
    """Create fact table with foreign keys"""

    # Join with dimensions to get surrogate keys
    fact_df = (raw_df
        .join(dim_entity, raw_df.entity_id ==
dim_entity.natural_key)
        .join(dim_date, raw_df.calculation_date ==
dim_date.calculation_date)
        .join(dim_product, raw_df.product_type ==
dim_product.product_type)
        .select(
            F.col("dim_entity.entity_key"),
            F.col("dim_date.date_key"),
            F.col("dim_product.product_key"),
            F.col("raw_df.liquidity_value"),
            F.col("raw_df.haircut_value"),
            F.col("raw_df.adjustment_flag"),
            F.col("raw_df.lineage_key"),
            F.current_timestamp().alias("load_timestamp")
        )
    )

    return fact_df

```

Redshift Table Design for Direct Write

-- Optimized fact table for direct writes

```

CREATE TABLE fact_liquidity (
    entity_key INTEGER NOT NULL ENCODE az64,
    date_key INTEGER NOT NULL ENCODE az64,
    product_key INTEGER NOT NULL ENCODE az64,
    liquidity_value DECIMAL(18,4) ENCODE az64,
    haircut_value DECIMAL(18,4) ENCODE az64,
    adjustment_flag BOOLEAN ENCODE raw,
    lineage_key VARCHAR(100) ENCODE zstd,
    load_timestamp TIMESTAMP ENCODE az64,

```

```

        PRIMARY KEY (entity_key, date_key, product_key)
    )
    DISTSTYLE KEY
    DISTKEY (entity_key)
    SORTKEY (date_key, entity_key)
    -- Optimize for concurrent writes
    ALTER TABLE fact_liquidity SET (
        append_only = true,
        backup = yes
    );

    -- Create staging table for efficient merges
    CREATE TABLE fact_liquidity_staging (LIKE fact_liquidity);

    -- Optimize for PowerBI queries
    CREATE MATERIALIZED VIEW mv_liquidity_current AS
    SELECT
        e.entity_name,
        d.calculation_date,
        p.product_category,
        SUM(f.liquidity_value) as total_liquidity,
        SUM(CASE WHEN f.adjustment_flag THEN f.liquidity_value ELSE 0 END)
            as adjusted_amount
    FROM fact_liquidity f
    JOIN dim_entity e ON f.entity_key = e.entity_key
    JOIN dim_date d ON f.date_key = d.date_key
    JOIN dim_product p ON f.product_key = p.product_key
    WHERE d.calculation_date >= DATEADD(day, -90, CURRENT_DATE)
    GROUP BY 1,2,3;

```

Performance Monitoring for Direct Writes

```

class DirectWriteMonitor:
    def __init__(self):
        self.metrics = {}

    def monitor_write_performance(self, write_func):
        def monitored_write(df, table_name):
            metrics = {
                "start_time": time.time(),
                "row_count": df.count(),
                "partition_count": df.rdd.getNumPartitions(),
                "table_name": table_name
            }

            # Monitor Redshift during write
            with self.monitor_redshift_load():
                result = write_func(df, table_name)

            metrics["duration"] = time.time() - metrics["start_time"]

```

```

        metrics["throughput"] = metrics["row_count"] /
        metrics["duration"]

        # Alert if performance degrades
        if metrics["throughput"] < self.config.min_throughput:
            self.alert_performance_degradation(metrics)

        return result

    return monitored_write

# Optimal S3 file structure
s3_optimization_config = {
    "file_format": "parquet",
    "compression": "zstd",
    "target_file_size": "128MB",
    "partition_strategy": "date_entity_type",
    "sort_keys": ["calculation_date", "entity_id", "metric_type"]
}

# Write optimization
def write_to_s3_optimized(df, path):
    (df.repartition(calculate_optimal_partitions(df))
     .sortWithinPartitions("calculation_date", "entity_id")
     .write
     .mode("overwrite")
     .option("compression", "zstd")
     .option("maxRecordsPerFile", 1000000)
     .parquet(path))

```

Redshift Optimization

```

-- Create optimized table structure
CREATE TABLE liquidity_metrics (
    calculation_date DATE ENCODE az64,
    entity_id VARCHAR(50) ENCODE zstd,
    metric_type VARCHAR(30) ENCODE bytedict,
    metric_value DECIMAL(18,4) ENCODE az64,
    adjustment_flag BOOLEAN ENCODE raw,
    lineage_key VARCHAR(100) ENCODE zstd
)
DISTSTYLE KEY
DISTKEY (entity_id)
SORTKEY (calculation_date, entity_id);

-- Materialized view for PowerBI
CREATE MATERIALIZED VIEW mv_liquidity_summary AS
SELECT
    calculation_date,
    entity_id,
    metric_type,

```

```

SUM(metric_value) as total_value,
MAX(CASE WHEN adjustment_flag THEN 1 ELSE 0 END) as
  has_adjustments,
COUNT(DISTINCT lineage_key) as calculation_paths
FROM liquidity_metrics
GROUP BY 1, 2, 3;

```

```
-- Auto-refresh strategy
```

```

ALTER MATERIALIZED VIEW mv_liquidity_summary
SET (AUTO REFRESH = YES);

```

PowerBI Optimization Strategies




Diagram 7

PowerBI Implementation Guidelines

1. Dataset Design

```

-- Create calculation groups for adjustments
CALCULATION GROUP 'Adjustment Scenarios'
  CALCULATION ITEM "Base" = SELECTEDMEASURE()
  CALCULATION ITEM "With Adjustments" =
    CALCULATE(
      SELECTEDMEASURE(),
      'Metrics'[adjustment_flag] = TRUE
    )
  CALCULATION ITEM "Variance" =
    [With Adjustments] - [Base]

```

2. Incremental Refresh Policy

```

{
  "incrementalRefresh": {
    "enabled": true,
    "rangeStart": "3 months",
    "rangeEnd": "0 days",
    "incrementalGranularity": "day",
    "detectDataChanges": "lastModified"
  }
}

```

```
}
}
```

3. Aggregation Strategy

```
-- Pre-aggregate in Redshift for common queries
CREATE TABLE agg_daily_liquidity AS
SELECT
    DATE_TRUNC('day', calculation_date) as day,
    entity_type,
    COUNT(DISTINCT entity_id) as entity_count,
    SUM(metric_value) as total_liquidity,
    AVG(metric_value) as avg_liquidity
FROM liquidity_metrics
GROUP BY 1, 2;
```

Implementation Patterns

Adjustment Service Implementation

```
class AdjustmentService:
    def __init__(self, spark_session, dependency_graph,
                 cache_manager):
        self.spark = spark_session
        self.dep_graph = dependency_graph
        self.cache = cache_manager

    def apply_adjustment(self, adjustment_rule):
        # Phase 1: Impact Analysis
        impact_analysis = self.analyze_impact(adjustment_rule)

        # Phase 2: Create execution plan
        exec_plan = self.create_execution_plan(impact_analysis)

        # Phase 3: Execute incremental updates
        results = self.execute_incremental_updates(exec_plan)

        # Phase 4: Update downstream systems
        self.propagate_changes(results)

        return AdjustmentResult(
            rule_id=adjustment_rule['rule_id'],
            affected_records=impact_analysis.affected_count,
            processing_time=results.duration,
            stages_updated=len(exec_plan.stages)
        )

    def execute_incremental_updates(self, exec_plan):
        results = []
```

```

for stage in exec_plan.stages:
    if stage.can_use_incremental:
        result = self.incremental_update(stage)
    else:
        result = self.full_recalculation(stage)

    results.append(result)

    # Update cache
    self.cache.update_stage(stage.id, result.data)

return ExecutionResults(results)

```

Lineage Tracking Implementation

```

class LineageTracker:
    def __init__(self):
        self.lineage_store = LineageStore()

    def wrap_transformation(self, transformation_func):
        def wrapped(df):
            # Extract input PKs
            input_pks = df.select("pk").distinct().collect()

            # Apply transformation
            result_df = transformation_func(df)

            # Extract output PKs
            output_pks = result_df.select("pk").distinct().collect()

            # Store lineage
            self.lineage_store.add_transformation(
                input_pks=input_pks,
                output_pks=output_pks,
                transformation_id=transformation_func.__name__,
                timestamp=datetime.now()
            )

            return result_df

        return wrapped

```

Monitoring and Observability

DAG Execution Monitoring

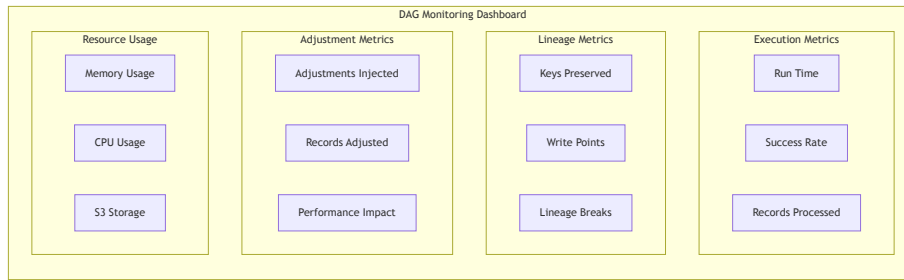


Diagram 8

Monitoring Implementation

```
class DAGMonitor:
    """
    Comprehensive monitoring for DAG-based pipelines
    """

    def __init__(self):
        self.metrics_collector = MetricsCollector()

    def monitor_operator_execution(self, operator: BaseOperator):
        """Monitor individual operator execution"""

        @functools.wraps(operator.execute)
        def monitored_execute(input_df):
            start_time = time.time()
            initial_keys = self.extract_keys(input_df)

            # Execute operator
            result_df = operator.execute(input_df)

            # Collect metrics
            metrics = {
                "operator_id": operator.operator_id,
                "operator_type": operator.__class__.__name__,
                "duration": time.time() - start_time,
                "input_records": input_df.count(),
                "output_records": result_df.count(),
                "keys_preserved": self.verify_key_preservation(
                    initial_keys,
                    self.extract_keys(result_df)
                ),
                "memory_usage": self.get_memory_usage(),
                "data_persisted": operator.should_persist()
            }
            return metrics
```



```

        self.metrics_collector.record(metrics)

        # Alert on anomalies
        if not metrics["keys_preserved"] and
operator.maintains_keys:
            self.alert_lineage_break(operator, metrics)

    return result_df

return monitored_execute

def monitor_adjustment_impact(self, adjustment_operator:
AdjustmentOperator):
    """Special monitoring for adjustment operators"""

    def monitored_adjustment(input_df):
        # Track before state
        before_snapshot = self.create_data_snapshot(input_df)

        # Apply adjustment
        result_df = adjustment_operator.execute(input_df)

        # Track after state
        after_snapshot = self.create_data_snapshot(result_df)

        # Calculate impact
        impact_metrics = {
            "records_modified": self.count_modified_records(
                before_snapshot,
                after_snapshot
            ),
            "value_change_distribution":
self.calculate_value_changes(
                before_snapshot,
                after_snapshot
            ),
            "affected_entities": self.get_affected_entities(
                before_snapshot,
                after_snapshot
            )
        }

        self.publish_adjustment_metrics(adjustment_operator,
            impact_metrics)

    return result_df

return monitored_adjustment

class OperatorMetricsCollector:
    """

```

Collects and aggregates operator-level metrics

"""

```
def __init__(self):
    self.metrics_store = []

def get_operator_performance_summary(self, dag_run_id: str) ->
    DataFrame:
    """
    Generate performance summary for all operators in a DAG run
    """

    metrics_df = self.spark.createDataFrame(self.metrics_store)

    return metrics_df.filter(
        F.col("dag_run_id") == dag_run_id
    ).groupBy("operator_id", "operator_type").agg(
        F.avg("duration").alias("avg_duration"),
        F.sum("input_records").alias("total_input_records"),
        F.sum("output_records").alias("total_output_records"),
        F.avg("memory_usage").alias("avg_memory_usage"),
        F.sum(F.when(F.col("data_persisted"), 1).otherwise(0))
            .alias("persistence_count")
    ).orderBy("avg_duration", ascending=False)
```

Best Practices and Recommendations

1. DAG Design Principles

- **Modular Operators:** Each operator should have a single, clear responsibility
- **Explicit Persistence:** Only persist when lineage would be lost (aggregations)
- **Adjustment Points:** Design DAGs with natural adjustment injection points
- **Parallel-Friendly:** Structure DAGs to allow parallel scenario execution

2. Lineage Management

- **Always Preserve Keys:** Every operator must maintain or map primary keys
- **Checkpoint Strategy:** Write checkpoints before any operation that loses row-level detail
- **Lineage Documentation:** Store lineage metadata alongside data files
- **Recovery Paths:** Design for easy replay from any checkpoint

3. Performance Optimization

- **Smart Caching:** Cache within Spark when data will be reused in the DAG
- **Lazy Persistence:** Only write to S3 when necessary for lineage
- **Partition Alignment:** Ensure operators respect and maintain partitioning
- **Resource Allocation:** Size Spark executors based on operator requirements

4. Adjustment Best Practices

- **Rule Validation:** Validate adjustment rules before DAG execution
- **Impact Preview:** Run test adjustments on sample data first
- **Parallel Scenarios:** Use parallel DAG execution for what-if analysis
- **Audit Trail:** Maintain complete history of all adjustments applied

5. Monitoring and Alerting

- **Operator SLAs:** Set performance thresholds for each operator type
- **Lineage Verification:** Alert immediately on unexpected lineage breaks
- **Resource Monitoring:** Track memory and CPU per operator
- **Data Quality Checks:** Validate output at each DAG stage

Example: Complete DAG with Adjustments

```
# Example of building a complete regulatory calculation DAG
def build_regulatory_dag():
    """
    Build a complete DAG with multiple adjustment injection points
    """

    builder = VisualDAGBuilder()

    # Define the base DAG
    dag_config = {
        "operators": [
            {"id": "ingest", "type": "ingestion", "params": {"source":
"deals"}},
            {"id": "enrich", "type": "enrichment", "params": {"refs":
["entities", "products"]}},
            {"id": "calc_base", "type": "calculation", "params":
{"metric": "base_liquidity"}},
            {"id": "calc_haircut", "type": "calculation", "params":
{"metric": "haircut"}},
            {"id": "agg_entity", "type": "aggregation", "params":
{"by": ["entity_id", "date"]}},
            {"id": "calc_final", "type": "calculation", "params":
{"metric": "lcr"}},
            {"id": "agg_final", "type": "aggregation", "params":
{"by": ["date"]}},
            {"id": "star_schema", "type": "star_schema", "params":
{"target": "liquidity_facts"}}
        ],
        "connections": [
            ("ingest", "enrich"),
            ("enrich", "calc_base"),
            ("calc_base", "calc_haircut"),
            ("calc_haircut", "agg_entity"),
            ("agg_entity", "calc_final"),
            ("calc_final", "agg_final")
        ]
    }
```

```

        ("agg_final", "star_schema")
    ]
}

# Build base DAG
for op in dag_config["operators"]:
    builder.add_operator(OperatorConfig(
        operator_type=op["type"],
        operator_id=op["id"],
        parameters=op["params"],
        inputs=[],
        outputs=[]
    ))

    for from_id, to_id in dag_config["connections"]:
        builder.connect(from_id, to_id)

# Example: Inject adjustments at different points
# Adjustment 1: After enrichment (deal-level adjustments)
builder.inject_adjustment(
    after_operator_id="enrich",
    adjustment_rules=load_deal_adjustments()
)

# Adjustment 2: After entity aggregation (entity-level adjustments)
builder.inject_adjustment(
    after_operator_id="agg_entity",
    adjustment_rules=load_entity_adjustments()
)

return builder.build_airflow_dag()

# Run parallel scenarios
analyzer = ParallelAdjustmentAnalyzer(build_regulatory_dag())
impact_report = analyzer.analyze_adjustment_impact(
    adjustment_scenarios=[
        AdjustmentScenario("conservative", conservative_rules),
        AdjustmentScenario("baseline", baseline_rules),
        AdjustmentScenario("optimistic", optimistic_rules)
    ],
    injection_points=["enrich", "calc_base", "agg_entity"]
)

```

Conclusion

This enhanced architecture provides a robust, performant, and auditable system for regulatory calculations with flexible adjustment capabilities. The combination of incremental processing, intelligent caching, and optimized data flow from S3 through

Redshift to PowerBI ensures both performance and accuracy while maintaining complete visibility into all calculations and adjustments.

Appendix A: Generic Star Schema Operator

Overview

A reusable PySpark operator that automatically generates a star schema from a set of input tables and their relationships. This operator analyzes the table relationships, identifies facts and dimensions, and creates the appropriate star schema structure.

Core Implementation

```
from pyspark.sql import SparkSession, DataFrame
import pyspark.sql.functions as F
from typing import Dict, List, Tuple, Set, Optional
from dataclasses import dataclass
from collections import defaultdict
import networkx as nx

@dataclass
class TableRelationship:
    """Defines a relationship between two tables"""
    from_table: str
    from_column: str
    to_table: str
    to_column: str
    relationship_type: str # "1:1", "1:n", "n:1", "n:n"

@dataclass
class ColumnMetadata:
    """Metadata about a column for schema generation"""
    column_name: str
    data_type: str
    is_measure: bool
    is_dimension: bool
    aggregation_func: Optional[str] = None # sum, avg, count, etc.

class StarSchemaOperator:
    """
    Generic Star Schema Generator for PySpark

    This operator takes a set of tables and their relationships and
    automatically
    generates a star schema by:
    1. Analyzing relationships to identify fact and dimension tables
    2. Creating surrogate keys for dimensions
    3. Building the fact table with proper foreign keys
    4. Handling slowly changing dimensions (SCD Type 2)
    """
```

```

"""

```

```

def __init__(self, spark: SparkSession):
    self.spark = spark
    self.dimension_tables = {}
    self.fact_tables = {}
    self.surrogate_key_mappings = {}

def generate_star_schema(
    self,
    tables: Dict[str, DataFrame],
    relationships: List[TableRelationship],
    column_metadata: Dict[str, List[ColumnMetadata]],
    target_fact_table: Optional[str] = None
) -> Dict[str, DataFrame]:
    """
    Main method to generate star schema

    Args:
        tables: Dictionary of table_name -> DataFrame
        relationships: List of relationships between tables
        column_metadata: Metadata about columns (measures vs
            dimensions)
        target_fact_table: Optional hint for main fact table

    Returns:
        Dictionary with generated star schema tables
    """

    # Step 1: Build relationship graph
    rel_graph = self._build_relationship_graph(relationships)

    # Step 2: Identify fact and dimension tables
    fact_tables, dim_tables = self._classify_tables(
        tables, rel_graph, column_metadata, target_fact_table
    )

    # Step 3: Generate dimension tables with surrogate keys
    dimension_dfs = self._generate_dimensions(tables, dim_tables,
        column_metadata)

    # Step 4: Generate fact tables
    fact_dfs = self._generate_facts(
        tables, fact_tables, dimension_dfs, relationships,
        column_metadata
    )

    # Step 5: Create any necessary bridge tables for many-to-many
    # relationships
    bridge_dfs = self._generate_bridge_tables(relationships,
        dimension_dfs)

```

```

# Combine all results
result = {}
result.update({f"dim_{name}": df for name, df in
dimension_dfs.items()})
result.update({f"fact_{name}": df for name, df in
fact_dfs.items()})
result.update({f"bridge_{name}": df for name, df in
bridge_dfs.items()})

return result

def _build_relationship_graph(self, relationships:
List[TableRelationship]) -> nx.DiGraph:
    """Build a directed graph of table relationships"""

    G = nx.DiGraph()

    for rel in relationships:
        G.add_edge(
            rel.from_table,
            rel.to_table,
            from_column=rel.from_column,
            to_column=rel.to_column,
            rel_type=rel.relationship_type
        )

    return G

def _classify_tables(
    self,
    tables: Dict[str, DataFrame],
    rel_graph: nx.DiGraph,
    column_metadata: Dict[str, List[ColumnMetadata]],
    target_fact_table: Optional[str]
) -> Tuple[Set[str], Set[str]]:
    """Classify tables as facts or dimensions based on
    relationships and metadata"""

    fact_tables = set()
    dim_tables = set()

    # If target fact table is specified, use it
    if target_fact_table and target_fact_table in tables:
        fact_tables.add(target_fact_table)

    # Analyze each table
    for table_name, df in tables.items():
        # Count measures vs dimensions
        measures = sum(
            1 for col in column_metadata.get(table_name, [])
            if col.is_measure
        )

```

```

dimensions = sum(
    1 for col in column_metadata.get(table_name, [])
    if col.is_dimension
)

# Tables with many measures are likely facts
if measures > dimensions:
    fact_tables.add(table_name)
# Tables with mostly dimensions are dimension tables
elif dimensions > measures:
    dim_tables.add(table_name)
else:
    # Use graph analysis - tables with many incoming edges
    are facts
    in_degree = rel_graph.in_degree(table_name)
    out_degree = rel_graph.out_degree(table_name)

    if in_degree > out_degree:
        fact_tables.add(table_name)
    else:
        dim_tables.add(table_name)

return fact_tables, dim_tables

def _generate_dimensions(
    self,
    tables: Dict[str, DataFrame],
    dim_tables: Set[str],
    column_metadata: Dict[str, List[ColumnMetadata]]
) -> Dict[str, DataFrame]:
    """Generate dimension tables with surrogate keys"""

    dimension_dfs = {}

    for dim_name in dim_tables:
        if dim_name not in tables:
            continue

        df = tables[dim_name]

        # Add surrogate key
        df_with_sk = df.withColumn(
            f"{dim_name}_key",
            F.monotonically_increasing_id()
        )

        # Add SCD Type 2 columns
        df_with_scd = df_with_sk.withColumn(
            "valid_from",
            F.current_timestamp()
        ).withColumn(

```



```

        "valid_to",
        F.lit("9999-12-31").cast("timestamp")
    ).withColumn(
        "is_current",
        F.lit(True)
    )

    # Select only dimension columns
    dim_columns = [
        col.column_name for col in
column_metadata.get(dim_name, [])
        if col.is_dimension
    ]

    # Always include the surrogate key and natural key
    columns_to_select = [f"{dim_name}_key"] + dim_columns + [
        "valid_from", "valid_to", "is_current"
    ]

    # Filter columns that exist in the dataframe
    existing_columns = df_with_scd.columns
    final_columns = [col for col in columns_to_select if col
in existing_columns]

    dimension_dfs[dim_name] =
df_with_scd.select(*final_columns)

    return dimension_dfs

def _generate_facts(
    self,
    tables: Dict[str, DataFrame],
    fact_tables: Set[str],
    dimension_dfs: Dict[str, DataFrame],
    relationships: List[TableRelationship],
    column_metadata: Dict[str, List[ColumnMetadata]]
) -> Dict[str, DataFrame]:
    """Generate fact tables with foreign keys to dimensions"""

    fact_dfs = {}

    for fact_name in fact_tables:
        if fact_name not in tables:
            continue

        fact_df = tables[fact_name]

        # Join with each related dimension to get surrogate keys
        for rel in relationships:
            if rel.from_table == fact_name and rel.to_table in
dimension_dfs:
                dim_df = dimension_dfs[rel.to_table]

```

```

        # Join to get surrogate key
        fact_df = fact_df.join(
            dim_df.filter(F.col("is_current") ==
True).select(
                F.col(rel.to_column),
                F.col(f"{rel.to_table}_key")
            ),
            fact_df[rel.from_column] ==
dim_df[rel.to_column],
            "left"
        )

        # Select measures and foreign keys
        measure_columns = [
            col.column_name for col in
column_metadata.get(fact_name, [])
            if col.is_measure
        ]

        # Get all foreign keys
        foreign_keys = [
            f"{dim_name}_key" for dim_name in dimension_dfs.keys()
            if f"{dim_name}_key" in fact_df.columns
        ]

        # Add fact metadata
        fact_df = fact_df.withColumn(
            "load_timestamp",
            F.current_timestamp()
        )

        # Select final columns
        final_columns = foreign_keys + measure_columns +
["load_timestamp"]
        existing_columns = fact_df.columns
        columns_to_select = [col for col in final_columns if col
in existing_columns]

        fact_dfs[fact_name] = fact_df.select(*columns_to_select)

    return fact_dfs

def _generate_bridge_tables(
    self,
    relationships: List[TableRelationship],
    dimension_dfs: Dict[str, DataFrame]
) -> Dict[str, DataFrame]:
    """Generate bridge tables for many-to-many relationships"""

    bridge_dfs = {}

```

```

# Find many-to-many relationships
for rel in relationships:
    if rel.relationship_type == "n:n":
        bridge_name = f"{rel.from_table}_{rel.to_table}"

        if rel.from_table in dimension_dfs and rel.to_table in
dimension_dfs:
            # Create bridge table with both surrogate keys
            from_df = dimension_dfs[rel.from_table].select(
                F.col(f"{rel.from_table}_key"),
                F.col(rel.from_column)
            )
            to_df = dimension_dfs[rel.to_table].select(
                F.col(f"{rel.to_table}_key"),
                F.col(rel.to_column)
            )

            # This is a simplified bridge - in reality would
            # need the actual
            # many-to-many relationship data
            bridge_df = from_df.crossJoin(to_df).select(
                F.col(f"{rel.from_table}_key"),
                F.col(f"{rel.to_table}_key"),

F.current_timestamp().alias("created_timestamp")
            )

            bridge_dfs[bridge_name] = bridge_df

return bridge_dfs

# Usage Example
def example_usage(spark: SparkSession):
    """Example of using the Star Schema Operator"""

    # Create sample data
    deals_df = spark.createDataFrame([
        ("D001", "P001", "E001", 1000000, 0.05, "2024-01-15"),
        ("D002", "P002", "E001", 2000000, 0.03, "2024-01-15"),
        ("D003", "P001", "E002", 1500000, 0.04, "2024-01-16"),
    ], ["deal_id", "product_id", "entity_id", "amount", "rate",
        "date"])

    entities_df = spark.createDataFrame([
        ("E001", "Entity One", "Banking", "US"),
        ("E002", "Entity Two", "Insurance", "UK"),
    ], ["entity_id", "entity_name", "sector", "country"])

    products_df = spark.createDataFrame([
        ("P001", "Bond", "Fixed Income", "Low"),
        ("P002", "Loan", "Credit", "Medium"),
    ], ["product_id", "product_name", "product_type", "risk_level"])

```

```
# Define tables
tables = {
    "deals": deals_df,
    "entities": entities_df,
    "products": products_df
}

# Define relationships
relationships = [
    TableRelationship("deals", "entity_id", "entities",
        "entity_id", "n:1"),
    TableRelationship("deals", "product_id", "products",
        "product_id", "n:1"),
]

# Define column metadata
column_metadata = {
    "deals": [
        ColumnMetadata("deal_id", "string", False, True),
        ColumnMetadata("amount", "double", True, False, "sum"),
        ColumnMetadata("rate", "double", True, False, "avg"),
        ColumnMetadata("date", "string", False, True),
    ],
    "entities": [
        ColumnMetadata("entity_id", "string", False, True),
        ColumnMetadata("entity_name", "string", False, True),
        ColumnMetadata("sector", "string", False, True),
        ColumnMetadata("country", "string", False, True),
    ],
    "products": [
        ColumnMetadata("product_id", "string", False, True),
        ColumnMetadata("product_name", "string", False, True),
        ColumnMetadata("product_type", "string", False, True),
        ColumnMetadata("risk_level", "string", False, True),
    ]
}

# Generate star schema
operator = StarSchemaOperator(spark)
star_schema = operator.generate_star_schema(
    tables,
    relationships,
    column_metadata,
    target_fact_table="deals"
)

# Display results
for table_name, df in star_schema.items():
    print(f"\n{table_name}:")
    df.show()
```

```
return star_schema
```

```
# Advanced Features Extension
```

```
class AdvancedStarSchemaOperator(StarSchemaOperator):
```

```
    """Extended version with additional features"""
```

```
def __init__(self, spark: SparkSession):
```

```
    super().__init__(spark)
```

```
    self.optimization_rules = {}
```

```
def add_date_dimension(self, start_date: str, end_date: str) ->
    DataFrame:
```

```
    """Generate a complete date dimension"""
```

```
    return (self.spark.sql(f"""
```

```
        WITH date_range AS (
```

```
            SELECT explode(sequence(
                to_date('{start_date}'),
                to_date('{end_date}'),
                interval 1 day
            )) as date
        )
```

```
    SELECT
```

```
        row_number() OVER (ORDER BY date) as date_key,
        date as calendar_date,
        year(date) as year,
        quarter(date) as quarter,
        month(date) as month,
        day(date) as day,
        date_format(date, 'EEEE') as day_name,
        date_format(date, 'MMMM') as month_name,
        weekofyear(date) as week_of_year,
        CASE
```

```
            WHEN date_format(date, 'E') IN ('Sat', 'Sun')
            THEN true ELSE false
```

```
        END as is_weekend,
```

```
        CASE
```

```
            WHEN date = last_day(date)
            THEN true ELSE false
```

```
        END as is_month_end
```

```
    FROM date_range
```

```
    """))
```

```
def optimize_for_redshift(self, star_schema: Dict[str, DataFrame])
    -> Dict[str, DataFrame]:
```

```
    """Apply Redshift-specific optimizations"""
```

```
    optimized = {}
```

```
    for table_name, df in star_schema.items():
```

```
        if table_name.startswith("dim_"):
```

```
            # Small dimensions – prepare for DISTSTYLE ALL
```

```

        optimized[table_name] = df.coalesce(1)
    elif table_name.startswith("fact_"):
        # Facts – sort by common query patterns
        key_columns = [col for col in df.columns if
col.endswith("_key")]
        if key_columns:
            optimized[table_name] =
df.sortWithinPartitions(*key_columns)
        else:
            optimized[table_name] = df
    else:
        optimized[table_name] = df

return optimized

def generate_aggregation_tables(
    self,
    fact_df: DataFrame,
    aggregation_levels: List[List[str]]
) -> Dict[str, DataFrame]:
    """Generate pre-aggregated tables for PowerBI performance"""

    agg_tables = {}

    for level in aggregation_levels:
        agg_name = "_".join(level)

        # Identify measure columns
        measure_cols = [
            col for col in fact_df.columns
            if fact_df.schema[col].dataType.simpleString() in
['double', 'float', 'decimal']
        ]

        # Create aggregation
        agg_df = fact_df.groupBy(*level).agg(
            *[F.sum(col).alias(f"sum_{col}") for col in
measure_cols],
            *[F.avg(col).alias(f"avg_{col}") for col in
measure_cols],
            F.count("*").alias("record_count")
        )

        agg_tables[f"agg_{agg_name}"] = agg_df

    return agg_tables

```

Integration with Your Pipeline

```

class LiquidityStarSchemaGenerator:
    """Specific implementation for liquidity calculations pipeline"""

```

```

def __init__(self, spark: SparkSession):
    self.spark = spark
    self.operator = AdvancedStarSchemaOperator(spark)

def generate_liquidity_star_schema(self, s3_paths: Dict[str, str])
    -> None:
    """Generate star schema from S3 data layers"""

    # Load final adjusted data from S3
    adjusted_df = self.spark.read.parquet(s3_paths["adjusted"])
    entities_df =
    self.spark.read.parquet(s3_paths["reference_entities"])
    products_df =
    self.spark.read.parquet(s3_paths["reference_products"])

    # Define the schema structure
    tables = {
        "liquidity_calculations": adjusted_df,
        "entities": entities_df,
        "products": products_df
    }

    relationships = [
        TableRelationship(
            "liquidity_calculations", "entity_id",
            "entities", "entity_id", "n:1"
        ),
        TableRelationship(
            "liquidity_calculations", "product_id",
            "products", "product_id", "n:1"
        ),
    ]

    column_metadata = {
        "liquidity_calculations": [
            ColumnMetadata("calculation_id", "string", False,
            True),
            ColumnMetadata("entity_id", "string", False, True),
            ColumnMetadata("product_id", "string", False, True),
            ColumnMetadata("liquidity_value", "double", True,
            False, "sum"),
            ColumnMetadata("haircut_value", "double", True, False,
            "sum"),
            ColumnMetadata("adjusted_value", "double", True,
            False, "sum"),
            ColumnMetadata("calculation_date", "date", False,
            True),
        ],
        # ... additional metadata
    }

    # Generate star schema
    star_schema = self.operator.generate_star_schema(

```

```

        tables, relationships, column_metadata,
        target_fact_table="liquidity_calculations"
    )

    # Add date dimension
    star_schema["dim_date"] = self.operator.add_date_dimension(
        "2020-01-01", "2025-12-31"
    )

    # Optimize for Redshift
    optimized_schema =
    self.operator.optimize_for_redshift(star_schema)

    # Generate aggregation tables for PowerBI
    agg_tables = self.operator.generate_aggregation_tables(
        star_schema["fact_liquidity_calculations"],
        [
            ["date_key", "entity_key"],
            ["date_key", "product_key"],
            ["date_key"],
        ]
    )
    optimized_schema.update(agg_tables)

    # Write to Redshift
    self._write_to_redshift(optimized_schema)

def _write_to_redshift(self, schema_dict: Dict[str, DataFrame]):
    """Write star schema to Redshift with appropriate
    distribution"""

    for table_name, df in schema_dict.items():
        if table_name.startswith("dim_"):
            # Dimensions use DISTSTYLE ALL
            dist_style = "ALL"
        elif table_name.startswith("fact_"):
            # Facts use DISTSTYLE KEY
            dist_style = "KEY"
        else:
            # Aggregations use AUTO
            dist_style = "AUTO"

        # Write with appropriate settings
        self._optimized_redshift_write(df, table_name, dist_style)

```

Key Features of the Star Schema Operator

1. **Automatic Classification:** Analyzes tables to identify facts vs dimensions
2. **Surrogate Key Generation:** Creates monotonically increasing IDs
3. **SCD Type 2 Support:** Includes validity dates and current flags
4. **Bridge Table Generation:** Handles many-to-many relationships

5. **Optimization Options:** Redshift-specific and PowerBI-specific optimizations

6. **Flexible Metadata:** Configurable column classification and aggregation rules

This operator can be used throughout your pipeline to automatically generate appropriate star schemas at any stage, making it easy to create different analytical views of your data.